
MODULE *Raft*

This is the formal specification for the *Raft* consensus algorithm.

Copyright 2014 *Diego Ongaro*.

This work is licensed under the Creative Commons Attribution – 4.0

International License <https://creativecommons.org/licenses/by/4.0/>

EXTENDS *Naturals*, *FiniteSets*, *Sequences*, *TLC*, *Client*, *Server*, *Messages*

VARIABLE *allStates*

All variables; used for stuttering (asserting state hasn't changed).

vars \triangleq $\langle allStates, messages, clientVars, serverVars, stateVars, followerVars, candidateVars, leaderVars, log$

Init \triangleq
 \wedge *InitClientVars*
 \wedge *InitServerVars*
 \wedge *InitMessageVars*
 \wedge *allStates* = 0

Defines how the variables may transition.

Next \triangleq
 $\wedge \vee \exists i \in Server : Restart(i)$
 $\quad \wedge$ UNCHANGED $\langle clientVars \rangle$
 $\vee \exists i \in Server : TimeoutFollower(i)$
 $\quad \wedge$ UNCHANGED $\langle clientVars \rangle$
 $\vee \exists i \in Server : TimeoutCandidate(i)$
 $\quad \wedge$ UNCHANGED $\langle clientVars \rangle$
 $\vee \exists i, j \in Server : RequestPreVote(i, j)$
 $\quad \wedge$ UNCHANGED $\langle clientVars \rangle$
 $\vee \exists i, j \in Server : RequestVote(i, j)$
 $\quad \wedge$ UNCHANGED $\langle clientVars \rangle$
 $\vee \exists i \in Server : BecomeCandidate(i)$
 $\quad \wedge$ UNCHANGED $\langle clientVars \rangle$
 $\vee \exists i \in Server : BecomeLeader(i)$
 $\quad \wedge$ UNCHANGED $\langle clientVars \rangle$
 $\vee \exists i \in Client, j \in Server : OpenSession(i, j)$
 $\quad \wedge$ UNCHANGED $\langle serverVars, stateVars, followerVars, candidateVars, leaderVars, logVars \rangle$
 $\vee \exists i \in Client, j \in Server : CloseSession(i, j)$
 $\quad \wedge$ UNCHANGED $\langle serverVars, stateVars, followerVars, candidateVars, leaderVars, logVars \rangle$
 $\vee \exists i \in Client, j \in Server : ClientRequest(i, j)$
 $\quad \wedge$ UNCHANGED $\langle serverVars, stateVars, followerVars, candidateVars, leaderVars, logVars \rangle$
 $\vee \exists i \in Server : AdvanceCommitIndex(i)$

$$\begin{aligned}
& \wedge \text{UNCHANGED } \langle \text{clientVars} \rangle \\
\vee \exists i \in \text{Server} : & \text{ApplyEntry}(i) \\
& \wedge \text{UNCHANGED } \langle \text{clientVars} \rangle \\
\vee \exists i, j \in \text{Server} : & \text{AppendEntries}(i, j) \\
& \wedge \text{UNCHANGED } \langle \text{clientVars} \rangle \\
\vee \exists m \in \text{DOMAIN } \text{messages} : & \text{Receive}(m) \\
& \wedge \text{UNCHANGED } \langle \text{clientVars} \rangle \\
\vee \exists m \in \text{DOMAIN } \text{messages} : & \text{DuplicateMessage}(m) \\
& \wedge \text{UNCHANGED } \langle \text{serverVars}, \text{stateVars}, \text{followerVars}, \text{candidateVars}, \text{leaderVars}, \text{logVars}, \text{clientVars} \rangle \\
\vee \exists m \in \text{DOMAIN } \text{messages} : & \text{DropMessage}(m) \\
& \wedge \text{UNCHANGED } \langle \text{serverVars}, \text{stateVars}, \text{followerVars}, \text{candidateVars}, \text{leaderVars}, \text{logVars}, \text{clientVars} \rangle \\
& \wedge \text{allStates}' = \text{allStates} + 1
\end{aligned}$$

$$\begin{aligned}
\text{Inv} & \triangleq \\
& \wedge \exists s1, s2 \in \text{Server} : \\
& \quad \text{Len}(\text{log}[s1]) > 0 \wedge \text{Len}(\text{log}[s2]) > 0 \Rightarrow \exists l1 \in \text{DOMAIN } \text{log}[s1], l2 \in \text{DOMAIN } \text{log}[s2] : \\
& \quad \quad l1 \leq \text{commitIndex}[s1] \wedge l2 \leq \text{commitIndex}[s2] \Rightarrow \text{log}[s1][l1] = \text{log}[s2][l2]
\end{aligned}$$

The specification must start with the initial state and transition according to *Next*.

$$\text{Spec} \triangleq \text{Init} \wedge \square[\text{Next}]_{\text{vars}}$$